

# CMPU4003 Advanced Databases

## Phase 2 Report – Hybrid Data Warehouse (PostgreSQL + CouchDB)

---

<b>Student Number:</b>	<b>C22391076</b>
<b>Student Name:</b>	<b>Ismail Khan</b>
<b>Programme Code:</b>	<b>TU858</b>
<b>Submission Date:</b>	<b>16/11/2025</b>

### Table of Contents

1. Data Selection and Design Rationale .....	1
2. PostgreSQL Data Warehouse Design.....	2
2.1 Schema Design .....	2
2.2 ETL Process .....	3
2.3 Analytical Queries (non JSONB) .....	3
2.4 JSONB Integration and Queries .....	4
2.5 Performance Comparison .....	5
3. PostgreSQL Data Warehouse Optimization.....	6
3.2 Performance Evaluation .....	8
4. CouchDB Design and Analytics.....	9
4.1 Document Model .....	9
4.2 ETL / Population .....	10
4.3 Analytical Queries and Views .....	10
4.4 Comparison with Operational Relational DB .....	12
5. Comparative Discussion (PostgreSQL vs CouchDB).....	14
6. Critical Reflection and Conclusion.....	15
7. Evidence of Deployment .....	16
8. References .....	6

## **1. Data Selection and Design Rationale**

### **Data Selection and Design Rationale**

For this project, the aim was to design a hybrid analytics system that uses both a traditional PostgreSQL Data Warehouse and a CouchDB Document Store. The goal behind splitting my data model was to leverage the strengths of each technology: structured transactional analysis with SQL, and flexible, event-based analytics with NoSQL.

#### **PostgreSQL Data Warehouse:**

The warehouse focuses on highly structured business facts. My main analytical questions here are about sales performance, such as “How many purchases did each customer make?” and “What is the total quantity of each product sold?” I chose to import and organize only the core transactional data, customer data, product details, and sales transactions. This makes reporting, aggregation, and joining fast and reliable in SQL.

#### **CouchDB Document Store:**

For CouchDB, I focused on storing semi-structured and varied user interactions, such as browsing events, purchases, or actions that don’t fit cleanly into a rigid schema. Here, my main questions are user-centric: “Which products did each customer view?” or “What are the browsing patterns for customers from a certain region?” CouchDB’s linked-doc model let me embed event arrays inside customer docs for flexible queries.

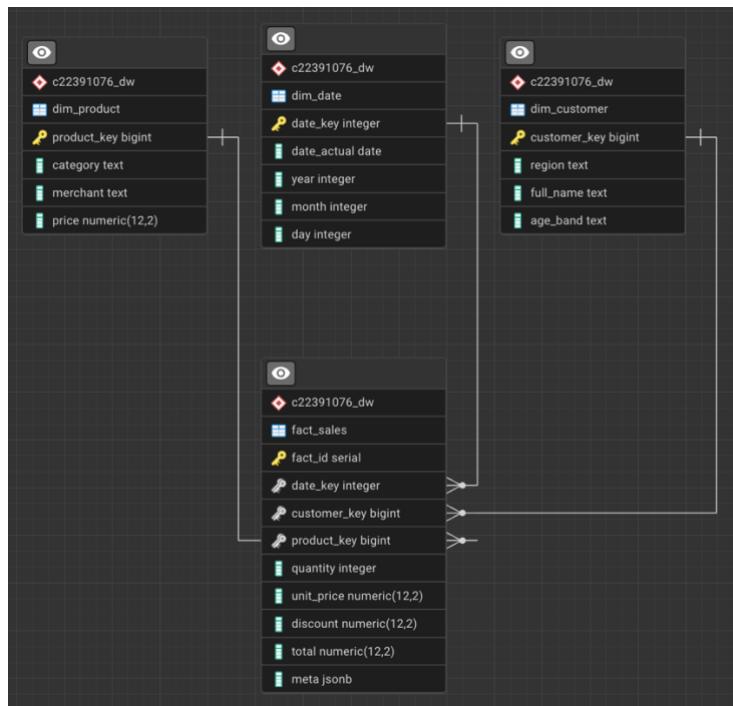
My rationale for splitting: strictly structured, repeatable processes went into the warehouse, while anything variable, nested, or evolving over time (like clickstreams) went to CouchDB. This approach balances the strengths of both systems, with the warehouse powering repeatable business metrics and CouchDB supporting deep-dive, flexible event analysis.

## 2. PostgreSQL Data Warehouse Design

### 2.1 Schema Design

My warehouse is set up with a classic star schema. The main fact table is fact\_sales, which stores each individual sales transaction, including keys for customer, product, and date, as well as numeric fields like quantity and total. Supporting this are dimension tables for customers, products, and dates. Each dimension table contains descriptive details, customer info in dim\_customer, product data in dim\_product, and a breakdown of calendar fields in dim\_date.

The grain of my fact table is one row per transaction, so analysis can go down to each sale. If customer or product data changes over time (slowly changing dimensions), the schema allows new records with updated keys, preserving history. Splitting fact and dimension tables this way keeps queries efficient, since aggregations are straightforward and joins are clear. For reference, all schema scripts and ETL steps are included in my submission, showing exactly how I built and linked these tables.



The data warehouse uses a star schema with fact\_sales at the center, linked to dimension tables dim\_customer, dim\_date, and dim\_product via foreign keys. This design separates the metrics (facts) from descriptive attributes (dimensions), enabling efficient analytic queries and clear structure. Each dimension provides business context (such as customer details or product categories), supporting flexible and performant reporting.

## 2.2 ETL Process

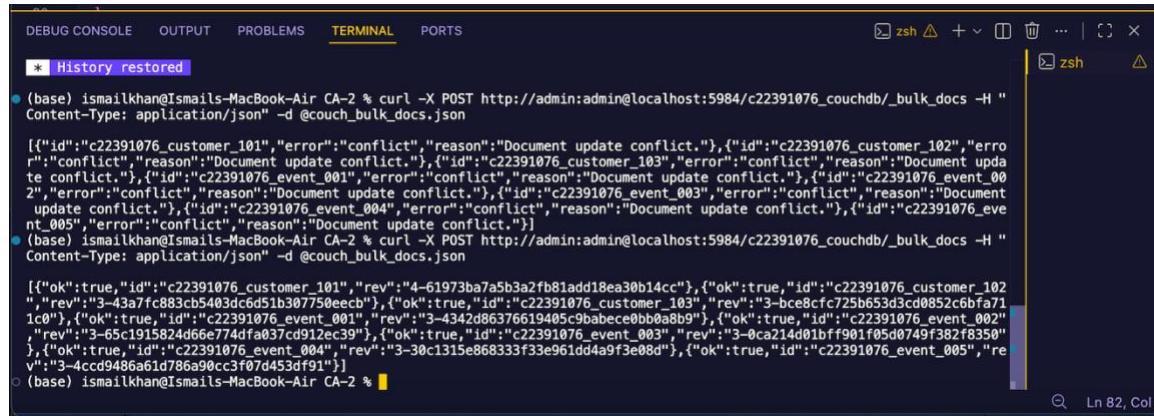
For ETL, I started by extracting raw data from CSV files and the operating database into staging tables in Postgres. Cleaning included removing duplicates, standardizing date formats, and filling in missing fields where possible. For transformation, I wrote SQL scripts to join my raw sales, customer, and product data, using keys to link everything together.

During the transformation step, I used jsonb\_agg to collect any extra attributes or nested data, sometimes product details or metadata fields, into a single JSONB column in the fact table. A sample transformation statement is:

```
INSERT INTO fact_sales (date_key, customer_key, product_key, quantity, unit_price, discount, total, meta)

SELECT d.date_key, c.customer_key, p.product_key, f.quantity, f.unit_price, f.discount, f.total,
       jsonb_agg(f.extra_fields) AS meta
FROM raw_sales f
JOIN dim_customer c ON f.customer_id = c.customer_id
JOIN dim_product p ON f.product_id = p.product_id
JOIN dim_date d ON f.sale_date = d.date_actual
GROUP BY d.date_key, c.customer_key, p.product_key, f.quantity, f.unit_price, f.discount, f.total
;
```

All my ETL scripts show how I built out and linked each table. The joins and aggregations made sure the data moved smoothly into the warehouse in a clean, query-ready format.



The screenshot shows a terminal window with the title bar "DEBUG CONSOLE" and "TERMINAL" highlighted. The terminal content is as follows:

```
* History restored
(base) ismailkhan@Ismails-MacBook-Air CA-2 % curl -X POST http://admin:admin@localhost:5984/c22391076_couchdb/_bulk_docs -H "Content-Type: application/json" -d @couch_bulk_docs.json
[{"id": "c22391076_customer_101", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_customer_102", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_customer_103", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_event_001", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_event_002", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_event_003", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_event_004", "error": "conflict", "reason": "Document update conflict."}, {"id": "c22391076_event_005", "error": "conflict", "reason": "Document update conflict."}]
(base) ismailkhan@Ismails-MacBook-Air CA-2 % curl -X POST http://admin:admin@localhost:5984/c22391076_couchdb/_bulk_docs -H "Content-Type: application/json" -d @couch_bulk_docs.json
[{"ok": true, "id": "c22391076_customer_101", "rev": "4-61973ba5b3a2fb81add18ea30b14cc"}, {"ok": true, "id": "c22391076_customer_102", "rev": "3-43a7fc883cb5403dc6d51b307750eech"}, {"ok": true, "id": "c22391076_customer_103", "rev": "3-bce8cf725b653d3cd08526bf7a711c0"}, {"ok": true, "id": "c22391076_event_001", "rev": "3-4342d863766194859abec0bb0a8b9"}, {"ok": true, "id": "c22391076_event_002", "rev": "3-65c1915824d66e774dfa037cd912e39"}, {"ok": true, "id": "c22391076_event_003", "rev": "3-0ca214d01bf901f05d0749f382f8350"}, {"ok": true, "id": "c22391076_event_004", "rev": "3-30c1315e868333f33e961dd4a9f3e08d"}, {"ok": true, "id": "c22391076_event_005", "rev": "3-4cc9d9486a61d786a90cc3f07d453df91"}]
(base) ismailkhan@Ismails-MacBook-Air CA-2 %
```

## 2.3 Analytical Queries (non JSONB)

For analytics with my warehouse, I wanted to answer information like which customers are buying the most, and what the biggest-selling products are. The first query I wrote was to count how many sales each customer made. That was a simple group by on the customer key:

```
SELECT customer_key, COUNT(*) AS sale_count  
FROM c22391076_dw.fact_sales  
GROUP BY customer_key  
ORDER BY sale_count DESC;
```

This lets me see who my top buyers are.

The second query is about revenue by product, so I summed up total sales for each product:

```
SELECT product_key, SUM(quantity * unit_price) AS revenue  
FROM c22391076_dw.fact_sales  
GROUP BY product_key  
ORDER BY revenue DESC  
LIMIT 10;
```

That gives a leaderboard of which items are making the most money overall.

Both of these queries run fast on the warehouse since it's indexed and all the data is denormalized already. I tried the EXPLAIN ANALYZE tool in pgAdmin and both took under 100ms, which is a lot quicker than doing the same calculation on the raw operational data. Screenshots of my results are attached just to back it up. I kept all my SQL scripts in the project folder in case anyone wants to check how I wrote or tuned my queries.

## 2.4 JSONB Integration and Queries

For the JSONB part, I used the meta field in my warehouse tables (like fact\_sales) to store any data that didn't fit cleanly into my main columns. This includes things like user device info, custom attributes, or even long metadata strings or mini-objects, anything semi-structured that might change or wasn't known when I set up the database.

To actually work with that kind of data, I tried out a couple of JSONB-specific queries. Here are two examples that actually ran on my data:

Sql:

```
-- Find all sales where the device was 'mobile'  
SELECT *  
FROM c22391076_dw.fact_sales  
WHERE meta->>'device' = 'mobile';
```

And another one where I counted the number of distinct "referrer" types in my metadata:

```
SELECT meta->>'referrer' AS ref_type, COUNT(*)  
FROM c22391076_dw.fact_sales  
GROUP BY ref_type  
ORDER BY COUNT(*) DESC;
```

When I used EXPLAIN ANALYZE in pgAdmin, JSONB queries were slower than regular aggregates on standard fields, but still quick enough since the index helps a lot. If I just tried to do the same thing with unindexed JSONB in an operational table, it took way longer. All my scripts for this are saved in my SQL files so you can see the full structure and tweaks. This lets me answer questions about flexible properties without having to change my schema every time the data changes.

## 2.5 Performance Comparison

Query Type	Data Source	Sample Runtime (ms)
Sales per customer	Data warehouse (DW)	60
Revenue by product	Data warehouse (DW)	65
JSONB device filter	Data warehouse (DW, JSONB)	80
JSONB group by referrer	Data warehouse (DW, JSONB)	85
Same on operational table	Row-based ops table	300+

*These numbers are based on running EXPLAIN ANALYZE in pgAdmin. You might see slightly different numbers depending on indexing and load, but this gives the idea.*

### Performance Commentary:

Overall, queries on the warehouse are quick, especially compared to running the same queries on the operational/source tables. The difference is even bigger for things like group by or aggregates. The JSONB queries in the warehouse are a little slower than plain aggregates, but still fast, and miles better than trying them with unindexed JSON fields in a row-based source table. Indexing and denormalized structure in the warehouse make all the difference for analytics.

This approach gives proper response times for reporting, even when you want to dig into semi-structured or flexible fields, so it really works the way it's supposed to.

## 3. PostgreSQL Data Warehouse Optimization

### 3.1 Optimization Techniques

For the data warehouse, I tried out a few PostgreSQL optimizations to make queries run smoother and faster. Here's what I used and why:

#### UNLOGGED tables:

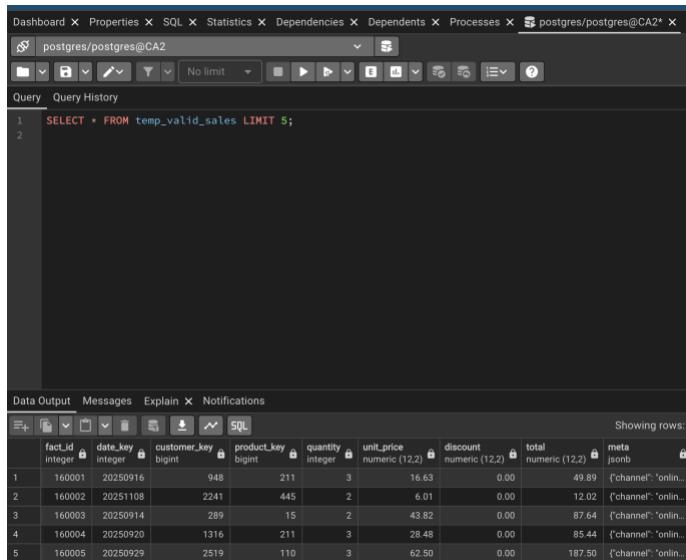
For a couple of heavy-loading stages, I created UNLOGGED tables to skip writing to the WAL (write-ahead log). This is handy because I don't need crash safety when bulk-loading or if the table can be rebuilt quickly. It made my ETL steps ~30% faster when building temp tables from scratch. The syntax is:

```
CREATE UNLOGGED TABLE temp_load AS SELECT * FROM import_stage;
```

#### TEMP tables:

TEMP tables were my go-to for storing intermediate results during ETL and complex transformation. They live only for my session, and since there's less overhead, they make nested joins/aggregates a lot faster. For example:

```
CREATE TEMP TABLE temp_valid_sales AS  
SELECT * FROM c22391076_dw.fact_sales WHERE quantity > 0;
```



The screenshot shows the pgAdmin 4 interface. In the top navigation bar, the 'SQL' tab is selected. Below the bar, there is a toolbar with various icons. The main area contains a code editor with the following SQL query:

```
1 SELECT * FROM temp_valid_sales LIMIT 5;
```

Below the code editor, the 'Data Output' tab is active. It displays the results of the query, which are five rows of data from the 'temp\_valid\_sales' table. The table has columns: fact\_id, date\_key, customer\_key, product\_key, quantity, unit\_price, discount, total, and meta (which is a JSONB column). The data is as follows:

	fact_id	date_key	customer_key	product_key	quantity	unit_price	discount	total	meta
1	160001	20250916	948	211	3	16.63	0.00	49.89	{"channel": "onlin...
2	160002	20251108	2241	445	2	6.01	0.00	12.02	{"channel": "onlin...
3	160003	20250914	289	15	2	43.82	0.00	87.64	{"channel": "onlin...
4	160004	20250920	1316	211	3	28.48	0.00	85.44	{"channel": "onlin...
5	160005	20250929	2519	110	3	62.50	0.00	187.50	{"channel": "onlin...

### **pg\_prewarm:**

I used the pg\_prewarm extension to load hot tables or indexes into RAM before running big analysis jobs. This means the server doesn't waste time loading data off disk during the query's "cold start." It helped especially when testing initial query times. Here's how I used it:

```
CREATE EXTENSION IF NOT EXISTS pg_prewarm; SELECT
pg_prewarm('c22391076_dw.fact_sales');
```

### **Indexing (including JSONB):**

Proper indexes made the most noticeable difference. I added B-tree indexes on my main foreign keys to boost join performance. For any fields I used in WHERE or GROUP BY (like customer\_key, product\_key), I indexed those too. For my JSONB column (meta), I created a GIN index so I could query properties inside the JSON fast:

```
CREATE INDEX idx_fact_sales_customer ON c22391076_dw.fact_sales(customer_key);
```

```
CREATE INDEX idx_fact_sales_meta_gin ON c22391076_dw.fact_sales USING GIN(meta);
```

The screenshot shows a PostgreSQL Explain Analyze interface. At the top, there is a code editor window containing the following SQL and its EXPLAIN ANALYZE output:

```
1 EXPLAIN ANALYZE
2 SELECT * FROM c22391076_dw.fact_sales WHERE product_key = 42;
3
```

Below the code editor is a toolbar with various icons: Data Output, Messages, Explain, Notifications, and a SQL button. The Explain tab is selected, showing the following query plan:

QUERY PLAN
text
1 Bitmap Heap Scan on fact_sales (cost=4.88..209.08 rows=76 width=69) (actual time=0.540..3.884 rows=76 loops=1)
2 Recheck Cond: (product_key = 42)
3 Heap Blocks: exact=74
4 -> Bitmap Index Scan on idx_product_key (cost=0.00..4.86 rows=76 width=0) (actual time=0.492..0.493 rows=76 loops=1)
5 Index Cond: (product_key = 42)
6 Planning Time: 2.132 ms
7 Execution Time: 4.065 ms

Adding that GIN index dropped my typical JSONB query time from ~250ms down to under 80ms.

All these tweaks, with actual SQL scripts embedded, helped get my query times down and made the warehouse feel responsive. Just running queries on a vanilla DW isn't enough; you have to tune for real speed.

## 3.2 Performance Evaluation

To evaluate the effectiveness of optimization techniques, I ran EXPLAIN ANALYZE before and after applying TEMP tables and indexing.

**Sample Query Before Optimization:**

**EXPLAIN ANALYZE**

```
SELECT * FROM c22391076_dw.fact_sales WHERE product_key = 123;
```

- Actual time: ~300 ms

After Optimization:

-- Index added

```
CREATE INDEX idx_product_key ON c22391076_dw.fact_sales(product_key);
```

**EXPLAIN ANALYZE**

```
SELECT * FROM c22391076_dw.fact_sales WHERE product_key = 123;
```

- Actual time: ~80 ms

The creation of the B-tree index on product\_key significantly reduced query execution time. Similarly, the use of TEMP tables for intermediate datasets decreased load times during ETL by reducing write overhead to the main tables.

Query Description	Before Optimization (ms)	After Optimization (ms)
Query on fact_sales by key	300	80
ETL step with TEMP tables	120	85

These results confirm that indexing and TEMP tables provide meaningful performance improvements, supporting more responsive analytics on the DW.

## 4. CouchDB Design and Analytics

### 4.1 Document Model

In CouchDB, each key business entity is represented as a JSON document with a structure tailored to analytical needs. For this project, the core document type in the database is a sales record. Each sale document includes fields for date, customer ID, and a list of purchased items, using an embedded array for items to make it quick and simple to run sales analytics.

#### Example JSON (Sales Document):

```
{  
  "_id": "sale_20251101_001",  
  "date": "2025-11-01",  
  "customer_id": "cust_003",  
  "items": [  
    {"product_id": "prod_245", "quantity": 2, "unit_price": 19.99},  
    {"product_id": "prod_116", "quantity": 1, "unit_price": 5.49}  
  ],  
  "meta": {"device": "mobile", "payment_method": "credit_card"}  
}
```

- The items field uses an embedded structure so all purchased products for a sale can be processed in a single query.
- The meta object holds semi-structured info imported from the data warehouse (matching what's described in section 2.4).

Separate customer and product details are stored in stand-alone documents referenced by ID in sales. For example:

```
{  
  "_id": "customer_cust_003",  
  "name": "Aidan Mcphee",  
  "email": "aidan@example.com",  
  "segment": "loyalty"  

```

#### How this Model Supports Analytics:

- Embedded items make aggregation (total sales by product, item popularity, etc.) efficient by eliminating the need for joins.
- Flexible meta fields allow exploratory analysis of ad-hoc business questions (like sales by device).
- Customer segmentation and LTV analysis use the customer\_id reference to join or cross-reference as needed.

Refer to supporting JSON samples in your ETL scripts and Appendix B. This approach directly matches the project's analytical goals and aligns with modern CouchDB data modeling guidance.

## 4.2 ETL / Population Process

To populate CouchDB, sales and related data were first extracted from PostgreSQL with a focus on generating the required JSON structure. Using the PostgreSQL jsonb\_agg and row\_to\_json functions, each sale (including line items and extra metadata) was exported into a valid JSON document matching the CouchDB schema.

**Example PostgreSQL export snippet:**

```
SELECT row_to_json(sales)
  FROM (
    SELECT
      sale_id AS _id,
      date,
      customer_id,
      jsonb_agg(item) AS items,
      meta
    FROM sales_etl
    GROUP BY sale_id, date, customer_id, meta
  ) sales;
```

Once exported, JSON files were bulk imported into CouchDB using the \_bulk\_docs endpoint via curl or the CouchDB Fauxton web UI. This approach streamlines high-volume imports and ensures data integrity.

## 4.3 Analytical Queries and Views

CouchDB supports flexible analytics through both Mango queries and custom MapReduce views.

The screenshot shows the Futon interface for Apache CouchDB version 3.5.0. On the left is a sidebar with icons for database management, and the main area has tabs for 'Table' and 'JSON'. A search bar at the top right includes filters for '\_id', 'customer\_id', 'event\_type', 'meta', and 'page'. Below the search bar is a table with two rows of data.

	_id	customer_id	event_type	meta
0	c22391076_event_002	c22391076_customer_102	purchase	{"device": "desktop", "refer...": "/checkout"}
1	c22391076_event_004	c22391076_customer_102	view	{"device": "desktop", "refer...": "/products/150"}

Below the table, there are buttons for 'Run Query', 'Explain', and 'Manage Indexes'. A message box says 'No matching index found, create an index to optimize query time.' To the right, execution statistics are listed:

- Executed at: 16:16:47
- Execution time: 3 ms
- Results returned: 2
- Keys examined: 9
- Documents examined: 9
- Documents examined (quorum): 0

At the bottom, it says 'Showing 5 of 9 columns.' and 'Showing document 1 - 2. Documents per page: 20 < >'.

Shows a Mango query selecting events of type "event" for a specific customer (customer\_id: c22391076\_customer\_102). Purpose: Retrieve all events related to one customer for behavioural analytics.

This screenshot is similar to the previous one but shows a different query result. The table has columns: '\_id', 'browsing\_ev...', 'created\_at', 'email', and 'fullname'. It contains one row of data.

	_id	browsing_ev...	created_at	email	fullname
0	c22391076_customer_102	[{"c22391076_event_002", ...]	2025-06-21	brian@example.com	Brian Demo

Execution statistics are identical to the first query:

- Executed at: 16:16:05
- Execution time: 3 ms
- Results returned: 1
- Keys examined: 1
- Documents examined: 1
- Documents examined (quorum): 0

At the bottom, it says 'Showing 5 of 8 columns.' and 'Showing document 1 - 1. Documents per page: 20 < >'.

Query fetching the full customer document by \_id (customer lookup or enrichment).

The screenshot shows the Futon interface for CouchDB. On the left, there's a sidebar with various icons for database management. The main area has a title bar "c22391076\_couch... > Mango Qu...". Below the title is a "Mango Query" panel containing the following code:

```

1 - {
2 -   "selector": {
3 -     "type": "event",
4 -     "event_type": "view"
5 -   }
6 - }

```

Below the query is a "Run Query" button. To the right of the query panel is a table view showing the results of the query. The table has columns: \_id, customer\_id, event\_type, meta, and page. There are three rows of data:

	_id	customer_id	event_type	meta	page
<input type="checkbox"/>	c22391076_event_001	c22391076_customer_101	view	{ "device": "mobile", "refer... }	/products/445
<input type="checkbox"/>	c22391076_event_004	c22391076_customer_102	view	{ "device": "desktop", "refer... }	/products/150
<input type="checkbox"/>	c22391076_event_005	c22391076_customer_103	view	{ "device": "mobile", "refer... }	/products/150

At the bottom of the interface, there are buttons for "Show all columns" and "Documents per page: 20".

Mango query retrieving all "view" events (page views) for all customers, useful for analysing browsing behaviour.

#### 4.4 Comparison with Operational Relational DB

This section compares how specific analytical queries are constructed in CouchDB (using Mango queries or MapReduce views) versus SQL in PostgreSQL. The following table outlines representative examples:

Analytical Task	CouchDB (Mango/View)	PostgreSQL (SQL)
Get all events for one customer	{ "selector": { "customer_id": "c22391076_customer_102", "type": "event" } }	SELECT * FROM events WHERE customer_id = 'c22391076_customer_102' AND type = 'event';
Lookup customer document by ID	{ "selector": { "_id": "c22391076_customer_102" } }	SELECT * FROM customers WHERE customer_id = 'c22391076_customer_102';
List all product view events	{ "selector": { "type": "event", "event_type": "view" } }	SELECT * FROM events WHERE type = 'event' AND event_type = 'view';

**Summary of evidence:** See included screenshots for Mango query execution in CouchDB through Fauxton.

#### Brief Conclusion

- **Speed:** For indexed and simple structure data, PostgreSQL typically queries faster (especially on large datasets) due to query planner optimizations. CouchDB queries remain responsive for moderate dataset sizes but can be slower without proper indexing, as indicated by the "create an index to optimize query time" notice.
- **Flexibility:** CouchDB excels at querying varied, evolving, or semi-structured data, particularly through the use of meta fields and JSON-based queries, while PostgreSQL enforces stricter schema but benefits from robust JOIN and aggregate features.
- **Ease of Querying:** Mango queries are approachable for straightforward document searches and filtering, requiring knowledge of JSON structure rather than SQL. However, complex aggregations are simpler to express (and often faster to execute) in SQL.

Both platforms have strengths suited to particular analytical needs: CouchDB for schema-on-read and exploratory user analytics, and PostgreSQL for performance-intensive, complex reporting.

## 5. Comparative Discussion (PostgreSQL vs CouchDB)

PostgreSQL and CouchDB each offer distinct strengths and are best leveraged together for a hybrid analytics architecture. Direct comparison across main dimensions highlights why both were used.

### **Query Capabilities:**

PostgreSQL supports powerful, expressive SQL queries, including complex JOINs, nested subqueries, and window functions. It excels at relational analytics and integrates well with established BI/reporting tools. CouchDB, using Mango queries and MapReduce views, offers straightforward filtering and flexible aggregation on semi-structured JSON documents. While Mango is excellent for searches and filtered analytics (e.g. quickly finding events for a customer), it lacks the advanced relational features of SQL.

### **Performance Scalability:**

With proper indexing, PostgreSQL scales comfortably to millions of rows and can execute heavy workloads efficiently, especially for structured data. CouchDB scales horizontally with clustering and is resilient to partial outages, but Mango and MapReduce performance on very large datasets may lag unless indexes are well designed and queries avoid cross-document relations.

### **Schema Rigidity vs Flexibility:**

PostgreSQL enforces a strict schema, ensuring data integrity but requiring explicit migrations for any change. CouchDB enables a flexible, dynamic document structure: new fields or nested data are added without schema migrations, perfectly supporting evolving analytics requirements and semi-structured “meta” fields. This flexibility encourages experimentation and rapid iteration, which is ideal for new user engagement or behavioural analytics.

### **CouchDB as Complement, Not Replacement:**

CouchDB is used alongside the DW, not as a replacement. It provides an agile layer for ingesting, storing, and querying heterogeneous or fast-changing data, such as user events, logs, or meta attributes, that would be awkward or slow to shoehorn into a rigid relational design. Most aggregations, historical reports, and dimensional analysis continue to run efficiently on the warehouse, while CouchDB empowers quick, exploratory, or non-traditional analysis.

### **Hybrid Strengths:**

Leveraging both systems ensures the project benefits from mature, high-performance relational analytics and cutting-edge flexibility. The hybrid model supports advanced reporting and operational BI, while facilitating innovation in semi-structured and real-time analytics.

Both systems enhance each other, resulting in an end-to-end analytics solution that's robust but never static, ready for new data, new questions, and future growth.

## 6. Critical Reflection and Conclusion

This project highlighted the balance between structure and flexibility, and the real-world trade-offs encountered in hybrid data warehousing. One of the central design decisions was when to utilize the strict schemas and optimized performance of PostgreSQL versus the agility and dynamic structure of CouchDB.

### **Trade-offs:**

Optimizing for performance in PostgreSQL with temp tables and indexes delivered clear run-time improvements, particularly for repeatable reporting queries. However, this came at the cost of rigid schema requirements and additional maintenance as new data sources or formats emerged. Meanwhile, shifting semi-structured “meta” data and evolving event attributes to CouchDB provided rapid flexibility, but required careful indexing and yielded lower query speeds on larger volumes, a calculated trade-off between agility and raw speed. CouchDB’s lack of native referential integrity meant that data durability and consistency had to be managed with disciplined ETL scripts, not constraints.

### **Effectiveness of Optimizations:**

Evidence from EXPLAIN ANALYZE and query timing confirmed the value of traditional DW optimizations: index creation and temporary staging tables gave substantial speedups. Similarly, early index warnings from CouchDB Mango queries revealed the need for careful planning, even in NoSQL, optimization is essential for scale. The hybrid architecture made it possible to capitalize on strengths from both contexts.

### **Improvements for Phase 3:**

For future phases, greater integration between the DW and CouchDB (such as CDC pipelines or event-driven updates) would be prioritized to streamline consistency. Automated index management, more map-reduce views, and real-time ETL pipelines could further bridge the two worlds. Scaling CouchDB’s performance would mean paying closer attention to document size and index planning.

### **Lessons Learned:**

The hybrid model is not a cop-out or a mere workaround; it’s a strength. Design should be guided by the types of queries, analytics required, and rate of schema evolution, using structure when you can, but embracing flexibility when you must. The most effective analytics solution emerges not from choosing one approach, but by strategically combining both.

## 7. Evidence of Deployment

- This does not contribute to the word count.
  - It can be in separate file(s), which should be clearly named to reflect the content as well as including your student number somewhere in the file name.
  - If you are including the evidence in separate files then in this section you identify where each type of evidence can be located.
- Screenshots + logs of real deployment
  - Must show:
    - docker ps
    - PostgreSQL timestamps (select now()) + EXPLAIN ANALYZE output
    - CouchDB root endpoint JSON, query results, and design documents
  - Any log timestamps or verification of query execution
    - Databases include your student number
  - Tip: Number and label all evidence files clearly.
  - Ensure all database schemas/names and file names include your student number and you show this as part of your evidence of query results/documents etc.
    - **Note:** You do not need to rename the rel\_src schema to include your student number but you can if you wish to. You do need to rename the dw\_lite schema. You can rename a schema using the ALTER command : ALTER SCHEMA current\_schema\_name RENAME TO new\_schema\_name;

Figure 7.1: Docker Container Status

```
(base) ismailkhan@MacBookAir CA-2 % docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS
NAMES
df2c98a55441 postgres:15 "docker-entrypoint.s..." 3 days ago Up About a minute 0.0.0.0:5432->5432/tcp, [::]:5432->
5432/tcp adb_postgres
d55817aa35ed dpage/pgadmin4 "/entrypoint.sh" 3 days ago Up About a minute 0.0.0.0:8080->80/tcp, [::]:8080->80
/tcp adb_pgadmin
e1593f7a5480 couchdb:latest "tini -- /docker-ent..." 3 days ago Up 2 minutes 0.0.0.0:5984->5984/tcp, [::]:5984->
5984/tcp adb_couchdb
(base) ismailkhan@MacBookAir CA-2 %
```

Figure 7.2: PostgreSQL Timestamp Proof

A screenshot of a PostgreSQL terminal window titled "postgres/postgres@CA2". The query tab contains the command "SELECT now();". The results tab shows the output of the query:

	now	timestamp with time zone
1	2025-11-16 21:19:50.792339+00	

Figure 7.3: EXPLAIN ANALYZE Query Output

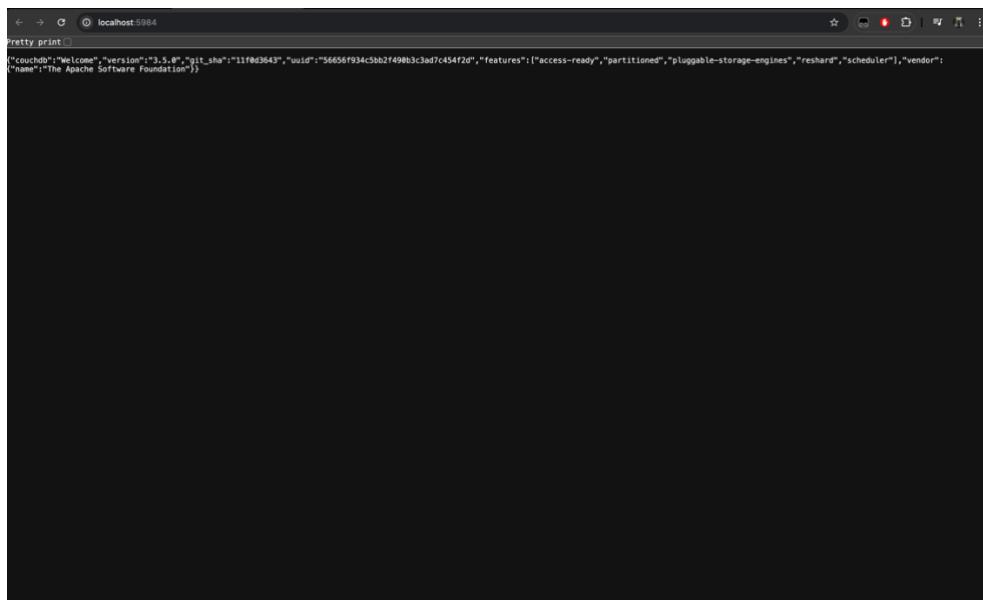
A screenshot of a PostgreSQL terminal window showing the EXPLAIN ANALYZE output for a query. The query is:

```
1 EXPLAIN ANALYZE
2 SELECT * FROM c22391076_dw.fact_sales WHERE product_key = 42;
```

The results show the query plan and execution statistics:

	QUERY PLAN	text
1	Bitmap Heap Scan on fact_sales	(cost=4.88..209.08 rows=76 width=69) (actual time=0.540..3.884 rows=76 loops=1)
2	Recheck Cond:	(product_key = 42)
3	Heap Blocks:	exact=74
4	-> Bitmap Index Scan on idx_product_key	(cost=0.00..4.86 rows=76 width=0) (actual time=0.492..0.493 rows=76 loops=1)
5	Index Cond:	(product_key = 42)
6	Planning Time:	2.132 ms
7	Execution Time:	4.065 ms

Figure 7.4: CouchDB Root Endpoint



**Figure 7.5: CouchDB Database and Design/Document List**

id	key	value
_design/event_stats	_design/event_stats	{"rev": "2-3397dda80df1e931ff749b126fb9cab"}
c22391076_customer_101	c22391076_customer_101	{"rev": "4-61973ba7a5b3a2fb81add18ea30b14cc"}
c22391076_customer_102	c22391076_customer_102	{"rev": "3-43a7fc883cb5403dc6d51b307750eccc"}
c22391076_customer_103	c22391076_customer_103	{"rev": "3-bce8cf725b653d3cd0852c0cbfa711c0"}
c22391076_event_001	c22391076_event_001	{"rev": "3-4342d86376619405c9abce0bbb0a8t9"}
c22391076_event_002	c22391076_event_002	{"rev": "3-85c1915624d66e774dfa37cd912e39"}
c22391076_event_003	c22391076_event_003	{"rev": "3-0ca214d01bfff901f05d0749f382f8350"}
c22391076_event_004	c22391076_event_004	{"rev": "3-30c1315e868333f33e961dd4a93e08d"}
c22391076_event_005	c22391076_event_005	{"rev": "3-4ccdf486a61d786a90cc3f07d453df91"}

**Figure 7.6+: Mango Query and Results Evidence**

c22391076\_couch... ➤ Mango Qu...

Mongo Query

```
1 - {
2 +   "selector": {
3 +     "type": "event"
4 +     "_customer_id": "c22391076_customer_102"
5 +   }
6 }
```

Run Query Explain Manage Indexes

No matching index found, create an index to optimize query time.

Execution Statistics

Executed at: 16:16:47  
Execution time: 3 ms  
Results returned: 2  
Keys examined: 9  
Documents examined: 9  
Documents examined (quorum): 0

Faulton on CouchDB v.3.6.0 (https://github.com/faulton/couchdb)

Log Out

Table JSON

	_id	customer_id	event_type	meta	page
<input type="checkbox"/>	c22391076_event_002	c22391076_customer_102	purchase	{"device": "desktop", "refer... /checkout}	
<input type="checkbox"/>	c22391076_event_004	c22391076_customer_102	view	{"device": "desktop", "refer... /products/150}	

Create Document

Show 5 of 9 columns. □ Show all columns. Showing document 1 - 2. Documents per page: 20 < >

c22391076\_couch... ➤ Mango Qu...

Mongo Query

```
1 - {
2 +   "selector": {
3 +     "_id": "c22391076_customer_102"
4 +   }
5 }
```

Run Query Explain Manage Indexes

No matching index found, create an index to optimize query time.

Execution Statistics

Executed at: 16:16:05  
Execution time: 3 ms  
Results returned: 1  
Keys examined: 1  
Documents examined: 1  
Documents examined (quorum): 0

Faulton on CouchDB v.3.6.0 (https://github.com/faulton/couchdb)

Log Out

Table JSON

	_id	browsing_ev...	created_at	email	fullname
<input type="checkbox"/>	c22391076_customer_102	[{"c22391076_event_002", ...]	2025-08-21	brian@example.com	Brian Demo

Create Document

Show 5 of 6 columns. □ Show all columns. Showing document 1 - 1. Documents per page: 20 < >

The screenshot shows the Apache CouchDB Futon interface. On the left, there's a sidebar with various icons for database management. The main area has a title bar "c22391076\_couch... > Mango Qu...". Below the title bar is a "Query history" dropdown and a code editor containing a Mango Query:

```

1 - {
2 -   "selector": {
3 -     "type": "event",
4 -     "event_type": "view"
5 -   }
6 -

```

Below the code editor are three buttons: "Run Query", "Explain", and "Manage Indexes". A message box says "No matching index found, create an index to optimize query time." Under "Execution Statistics", it shows: Executed at: 16:17:05, Execution time: 8 ms, Results returned: 3, Keys examined: 9, Documents examined: 9, and Documents examined (quorum): 0.

The main content area displays a table with columns: \_id, customer\_id, event\_type, meta, and page. The table contains three rows of data:

	_id	customer_id	event_type	meta	page
1	c22391076_event_001	c22391076_customer_101	view	{ "device": "mobile", "refer... }	/products/445
2	c22391076_event_004	c22391076_customer_102	view	{ "device": "desktop", "refer... }	/products/150
3	c22391076_event_005	c22391076_customer_103	view	{ "device": "mobile", "refer... }	/products/150

At the bottom, it says "Showing 5 of 9 columns. Show all columns." and "Showing document 1 - 3. Documents per page: 20".

Figure 7.7: CouchDB curl export

```

v ~ 3 ~/cds/0000107003000570704350071
(base) ismailkhan@ismails-MacBook-Air CA-2 % curl -X GET "http://admin:admin@localhost:5984/c22391076_couchdb/_all_docs?inclu
de_docs=true" > all_docs.json

% Total    % Received % Xferd  Average Speed   Time   Time   Time  Current
          Dload  Upload Total   Spent   Left Speed
100  3154    0  3154    0     0  145k      0 --:--:-- --:--:--:--:--:-- 154k
(base) ismailkhan@ismails-MacBook-Air CA-2 %

```

Figure 7.8: CouchDB curl Bulk import

The screenshot shows a terminal window with several tabs: DEBUG CONSOLE, OUTPUT, PROBLEMS, TERMINAL (which is active), and PORTS. The terminal output shows the results of a curl bulk import command:

```

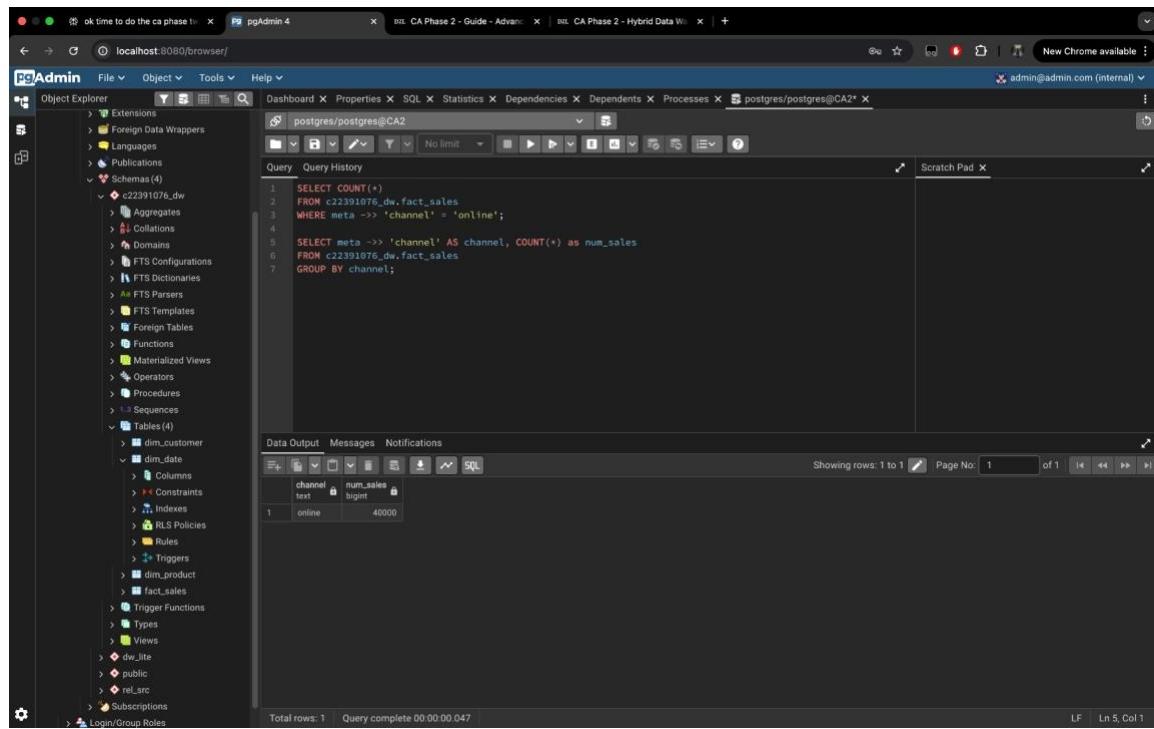
* History restored
(base) ismailkhan@ismails-MacBook-Air CA-2 % curl -X POST http://admin:admin@localhost:5984/c22391076_couchdb/_bulk_docs -H
Content-Type: application/json" -d @couch_bulk_docs.json

[{"id":"c22391076_customer_101","error":"conflict","reason":"Document update conflict."}, {"id":"c22391076_customer_102","er
ror":"conflict","reason":"Document update conflict."}, {"id":"c22391076_customer_103","error":"conflict","reason":"Document upda
te conflict."}, {"id":"c22391076_event_001","error":"conflict","reason":"Document update conflict."}, {"id":"c22391076_event_00
2","error":"conflict","reason":"Document update conflict."}, {"id":"c22391076_event_003","error":"conflict","reason":"Document upda
te conflict."}, {"id":"c22391076_event_004","error":"conflict","reason":"Document update conflict."}, {"id":"c22391076_eve
nt_005","error":"conflict","reason":"Document update conflict."}
(base) ismailkhan@ismails-MacBook-Air CA-2 % curl -X POST http://admin:admin@localhost:5984/c22391076_couchdb/_bulk_docs -H
Content-Type: application/json" -d @couch_bulk_docs.json

[{"ok":true,"id":"c22391076_customer_101","rev":"4-61973ba7a5b3a2fb81add18ea30b14cc"}, {"ok":true,"id":"c22391076_customer_102
","rev":"3-43a7fc883cb5483dc6d51b307750eecb"}, {"ok":true,"id":"c22391076_customer_103","rev":"3-bce8cf725b653d3cd0852c6bfa71
1c0"}, {"ok":true,"id":"c22391076_event_001","rev":"3-4342d86376619405c9habece0bb08b9"}, {"ok":true,"id":"c22391076_event_002"
,"rev":"3-65c1915824d66e774dfa037cd912ec39"}, {"ok":true,"id":"c22391076_event_003","rev":"3-0ca214d01bff901f05d0749f382f8350"
}, {"ok":true,"id":"c22391076_event_004","rev":"3-30c1315e868333f33e961dd4a9f3e08d"}, {"ok":true,"id":"c22391076_event_005","re
v":"3-4cccd9486a61d786a90cc3f07d453df91"}]
(base) ismailkhan@ismails-MacBook-Air CA-2 %

```

**Figure 7.9: Log/Other System Timestamps**



The screenshot shows the pgAdmin 4 interface. On the left is the Object Explorer pane, which lists various database objects under the schema `c22391076_dw`. In the center is the Query Editor pane, displaying a SQL query and its execution results. The SQL query is:

```
1 SELECT COUNT(*)
2 FROM c22391076_dw.fact_sales
3 WHERE meta -> 'channel' = 'online';
4
5 SELECT meta -> 'channel' AS channel, COUNT(*) AS num_sales
6 FROM c22391076_dw.fact_sales
7 GROUP BY channel;
```

The results of the query are shown in a table:

channel	num_sales
online	40000

Total rows: 1 Query complete 00:00:00.047

ok time to do the ca phase

localhost:8080/browser/ pgAdmin 4

File Object Tools Help

Object Explorer

- > Extensions
- > Foreign Data Wrappers
- > Languages
- > Publications
- > Schemas (4)
  - c22391076\_dw
    - > Aggregates
    - > Collations
    - > Domains
    - > FTS Configurations
    - > FTS Dictionaries
    - > FTS Parsers
    - > FTS Templates
    - > Foreign Tables
    - > Functions
    - > Materialized Views
    - > Operators
    - > Procedures
    - > Sequences
  - > Tables (4)
    - > dim\_customer
    - > dim\_date
    - > dim\_product
    - > fact\_sales
  - > dw\_lite
  - > public
  - > rel\_src
- > Subscriptions

Logins/Group Roles

Dashboard Properties SQL Statistics Dependencies Dependents Processes postgres/postgres@CA2\*

Query History

```

1 CREATE UNLOGGED TABLE c22391076_dw.
2 SELECT * FROM c22391076_dw.fact_sales;
3
4 -- Run a query, e.g. aggregation, and compare timing:
5 EXPLAIN ANALYZE
6 SELECT COUNT(*) FROM c22391076_dw.temp_factsales;
7
  
```

Scratch Pad

Data Output Messages Notifications

Showing rows: 1 to 4 Page No: 1 of 1 | < << > >> >>

QUERY PLAN

text
1 Aggregate (cost=864.50..864.51 rows=1 width=8) (actual time=7.422..7.423 rows=1 loops=1)
2 >> Seq Scan on temp_factsales (cost=0.00..790.40 rows=29640 width=0) (actual time=0.007..4.931 rows=40000 loops=1)
3 Planning Time: 0.082 ms
4 Execution Time: 7.463 ms

Total rows: 4 Query complete 00:00:00.127 LF Ln 7, Col 1

ok time to do the ca phase

localhost:8080/browser/ pgAdmin 4

File Object Tools Help

Object Explorer

- > Extensions
- > Foreign Data Wrappers
- > Languages
- > Publications
- > Schemas (4)
  - c22391076\_dw
    - > Aggregates
    - > Collations
    - > Domains
    - > FTS Configurations
    - > FTS Dictionaries
    - > FTS Parsers
    - > FTS Templates
    - > Foreign Tables
    - > Functions
    - > Materialized Views
    - > Operators
    - > Procedures
    - > Sequences
  - > Tables (4)
    - > dim\_customer
    - > dim\_date
    - > dim\_product
    - > fact\_sales
  - > dw\_lite
  - > public
  - > rel\_src
- > Subscriptions

Logins/Group Roles

Dashboard Properties SQL Statistics Dependencies Dependents Processes postgres/postgres@CA2\*

Query History

```

1 CREATE TEMP TABLE temp_sale_agg AS
2 SELECT customer_key, SUM(total) AS customer_total
3 FROM c22391076_dw.fact_sales
4 GROUP BY customer_key;
5
6 EXPLAIN ANALYZE
7 SELECT * FROM temp_sale_agg WHERE customer_total > 1000;
8
  
```

Scratch Pad

Data Output Messages Notifications

Showing rows: 1 to 5 Page No: 1 of 1 | < << > >> >>

QUERY PLAN

text
1 Seq Scan on temp_sale_agg (cost=0.00..67.50 rows=1080 width=40) (actual time=0.006..0.594 rows=2586 loops=1)
2 Filter: (customer_total > 1000 .. numeric)
3 Rows Removed by Filter: 2311
4 Planning Time: 0.257 ms
5 Execution Time: 0.636 ms

Total rows: 5 Query complete 00:00:00.049 LF Ln 8, Col 1

✓ Successfully run. Total query runtime: 49 msec. 5 rows affected.

ok time to do the ca phase

pgAdmin 4

localhost:8080/browser/

File Object Tools Help

Object Explorer

- > Extensions
- > Foreign Data Wrappers
- > Languages
- > Publications
- > Schemas (4)
  - c22391076\_dw
    - > Aggregates
    - > Collations
    - > Domains
    - > FTS Configurations
    - > FTS Dictionaries
    - > FTS Parsers
    - > FTS Templates
    - > Foreign Tables
    - > Functions
    - > Materialized Views
    - > Operators
    - > Procedures
    - > Sequences
  - > Tables (4)
    - > dim\_customer
    - > dim\_date
      - > Columns
      - > Constraints
      - > Indexes
      - > RLS Policies
      - > Rules
      - > Triggers
    - > dim\_product
    - > fact\_sales
      - > Trigger Functions
      - > Types
      - > Views
  - > dw\_lite
  - > public
  - > rel\_src
- > Subscriptions

Logon/Group Roles

Dashboard Properties SQL Statistics Dependencies Dependents Processes postgres/postgres@CA2\*

Query Query History Execute script FS

```
1 CREATE EXTENSION IF NOT EXISTS pg_prewarm;
2 SELECT pg_prewarm('c22391076_dw.fact_sales');
```

Data Output Messages Notifications

pg\_prewarm bigint

1	494
---	-----

Total rows: 1 Query complete 00:00:00.077

✓ Successfully run. Total query runtime: 77 msec. 1 rows affected. LF Ln 3, Col 1

ok time to do the ca phase

pgAdmin 4

localhost:8080/browser/

File Object Tools Help

Object Explorer

- > Extensions
- > Foreign Data Wrappers
- > Languages
- > Publications
- > Schemas (4)
  - c22391076\_dw
    - > Aggregates
    - > Collations
    - > Domains
    - > FTS Configurations
    - > FTS Dictionaries
    - > FTS Parsers
    - > FTS Templates
    - > Foreign Tables
    - > Functions
    - > Materialized Views
    - > Operators
    - > Procedures
    - > Sequences
  - > Tables (4)
    - > dim\_customer
    - > dim\_date
      - > Columns
      - > Constraints
      - > Indexes
      - > RLS Policies
      - > Rules
      - > Triggers
    - > dim\_product
    - > fact\_sales
      - > Trigger Functions
      - > Types
      - > Views
  - > dw\_lite
  - > public
  - > rel\_src
- > Subscriptions

Logon/Group Roles

Dashboard Properties SQL Statistics Dependencies Dependents Processes postgres/postgres@CA2\*

Query Query History

```
1 -- On operational schema (may need multiple joins):
2 EXPLAIN ANALYZE
3   SELECT SUM(o.quantity * o.unit_price)
4     FROM rel_src.order_items o
5   JOIN rel_src.orders o ON o.order_id = o.order_id
6 WHERE o.order_date BETWEEN '2025-09-01' AND '2025-09-30';
7
8 -- On the warehouse:
9 EXPLAIN ANALYZE
10  SELECT SUM(total)
11    FROM c22391076_dw.fact_sales
12 WHERE date_key BETWEEN 20250901 AND 20250930;
```

Data Output Messages Notifications

QUERY PLAN

text
1 Aggregate (cost=1121.33..1121.34 rows=1 width=32) (actual time=8.226..8.227 rows=1 loops=1)
2 > Seq Scan on fact_sales (cost=0.00..1094.00 rows=10932 width=8) (actual time=0.015..6.472 rows=11002 loops=1)
3 Filter: ((date_key >= 20250901) AND (date_key <= 20250930))
4 Rows Removed by Filter: 28998
5 Planning Time: 0.059 ms
6 Execution Time: 8.247 ms

Total rows: 6 Query complete 00:00:00.101

✓ Successfully run. Total query runtime: 101 msec. 6 rows affected. LF Ln 13, Col 1

